

# IMRA — Intelligent Multi-Agent Research Assistant

**Goal:** Build a polished, production-ready capstone demonstrating advanced agentic features (ADK + Gemini) that uses **>5 elements** learned during the 5-Day AI Agents Intensive. This project is designed to win — it showcases architecture, multi-agent collaboration (A2A), tools, MCP/long-running ops, memory, observability, evaluation, and optional Vertex AI deployment.

---

## Table of contents

1. Project summary (elevator pitch)
  2. Elements covered (mapping to Day 1–5)
  3. System architecture (components + data flow)
  4. Agent roles & definitions (detailed)
  5. Tooling & example code snippets (ADK-style Python skeletons)
  6. A2A message flows and MCP usage
  7. Memory design
  8. Observability plan (logs/traces/metrics)
  9. Evaluation strategy (LLM-as-a-Judge + metrics)
  10. Notebook & folder structure (Kaggle-ready)
  11. Deployment plan (optional — Vertex AI Agent Engine)
  12. Demo script + testing plan
  13. Slides / Presentation outline (10 slides)
  14. LinkedIn capstone post (short + carousel caption)
  15. Scoring rubric alignment (to maximize marks)
  16. Next steps & immediate deliverables
- 

## 1) Project summary (elevator pitch)

IMRA is a multi-agent research co-pilot that—given a research query—searches the web and PDFs, extracts key findings, synthesizes insights, evaluates quality, and returns a final summary plus confidence & references. It demonstrates:

- Multi-agent cooperation via A2A protocol
- Tool-calling (Search, PDF reader, chart generator)
- Session + long-term memory for user preferences
- MCP for long-running approvals and tool outputs
- Observability and LLM-as-a-Judge evaluation
- Optional deployment to Vertex AI Agent Engine

**User Story:** "Prepare a 500-word literature summary on 'micro-Doppler signatures for human activity recognition' with 5 citations and a short critique." IMRA returns a complete package: summary, short annotated bibliography, confidence scores, and visual charts.

---

## 2) Elements covered (choose >5)

1. Build Agent using Gemini + ADK (skeleton code provided)

2. Multi-Agent System (Research, Analysis, Critique, Memory agents)
  3. Tool Calling / Custom Tools (search, PDF scraper, chart maker)
  4. MCP (long-running operations + approvals)
  5. Sessions & Memory (short session + persistent memory)
  6. Observability (logs, traces, metrics)
  7. Evaluation (LLM-as-a-Judge, automated metrics)
  8. A2A Protocol (agent-to-agent messaging)
  9. Optional: Vertex AI Agent Engine deployment
- 

### 3) System architecture (components + data flow)

**Components:** - **User Interface:** Simple CLI or web form in notebook (input prompt + optional file uploads) - **Coordinator / Orchestrator Agent:** Receives user request, splits tasks, and orchestrates other agents - **Research Agent:** Runs Search + PDF tools to gather sources - **Analysis Agent:** Extracts, consolidates, generates summary and charts - **Critique Agent (Judge):** Evaluates outputs using LLM-as-a-Judge - **Memory Agent:** Stores user preferences & previous queries - **Tools Layer:** Search API, PDF reader, Chart generator, Citation formatter - **Observability Layer:** Centralized logging & traces (OpenTelemetry + local logs) - **Persistence:** SQLite or simple JSON store for demo (or Firestore/Redis if deployed)

**Data flow:** 1. User -> Orchestrator 2. Orchestrator -> Research Agent (A2A) 3. Research Agent -> Tools (search/pdf) -> returns sources 4. Research -> Analysis Agent (A2A) -> produce summary + charts 5. Analysis -> Critique Agent -> quality score & suggestions 6. Orchestrator collects final package -> returns to user 7. All steps emit logs/traces; results stored to memory

---

### 4) Agent roles & definitions (detailed)

#### Orchestrator Agent

- Responsibilities: Validate user request, allocate tasks, handle MCP confirmations, combine results.
- Inputs: user query, preferences from Memory Agent
- Outputs: final package (summary + citations + charts)

#### Research Agent

- Responsibilities: Use Search tool, fetch PDFs, extract key sentences, create bibliography candidates.
- Tools: Web Search tool, PDF extractor
- Output: List of candidate sources with short extracts

#### Analysis Agent

- Responsibilities: Merge extracts, deduplicate, generate final summary, create charts (word cloud, keyword frequency), create annotated bibliography
- Tools: summarizer (LLM), small python code for charts
- Output: Summary, charts, bibliography

## Critique Agent (LLM-as-a-Judge)

- Responsibilities: Evaluate summary on clarity, accuracy, citations, originality; give scores and suggested improvements
- Tools: A separate LLM prompt (can use the same Gemini model but with evaluation prompt template)
- Output: Scores & comments

## Memory Agent

- Responsibilities: store user preferences and previous sessions (e.g., preferred citation style)
- Storage: lightweight JSON DB or SQLite

## 5) Tooling & example code snippets

Below are ADK-like skeletons (Python). Replace `gemini_client` placeholders with your ADK/SDK calls.

### a) Tool: `web_search_tool.py` (skeleton)

```
# web_search_tool.py
class WebSearchTool:
    def __init__(self, api_key=None):
        self.api_key = api_key

    def search(self, query, top_k=5):
        # placeholder: call search API (Google Search API or SerpAPI)
        # return list of {title, url, snippet}
        results = []
        return results

    def fetch_page(self, url):
        # fetch HTML, extract main text
        return "...text..."
```

### b) Tool: `pdf_extractor.py`

```
import fitz # PyMuPDF

def extract_text_from_pdf(path, max_pages=5):
    text = ""
    doc = fitz.open(path)
    for p in doc:
        text += p.get_text()
        if len(text) > 20000:
            break
    return text
```

### c) Agent skeleton (ADK-like)

```
# orchestrator.py
from adk import Agent # pseudo

class OrchestratorAgent(Agent):
    def on_message(self, user_input, session):
        # validate, call memory, spawn research agent
        research_out = self.call_agent('research_agent', {'query': user_input})
        analysis_out = self.call_agent('analysis_agent', research_out)
        critique_out = self.call_agent('critique_agent', analysis_out)
        final = self.combine(research_out, analysis_out, critique_out)
        return final
```

Note: The ADK used in Kaggle codelabs uses a particular Python SDK. The pseudocode above maps to that structure. During notebook coding I'll give exact ADK call style.

## 6) A2A message flows and MCP usage

A2A messages (JSON payload):

```
{
  "from": "orchestrator",
  "to": "research_agent",
  "type": "task.request",
  "task": {"query": "...", "top_k": 5},
  "session_id": "sess-123"
}
```

**MCP long-running op example:** - Use MCP when a tool needs human approval (e.g., when fetching a paywalled PDF or performing a long scrape). - Orchestrator will create a long-running op with status PENDING\_APPROVAL and wait for user confirm.

Flow: 1. Orchestrator -> ResearchAgent: start long-running search 2. ResearchAgent returns op\_id and status=PENDING\_APPROVAL with sample excerpt 3. User confirms -> Orchestrator resumes ResearchAgent -> tool continues

## 7) Memory design

**Two-level memory:** - **Session memory (short-term):** In-memory dict stored per run; keeps conversation history - **Long-term memory:** SQLite/JSON table: sessions(id, user\_id, query, summary, timestamp, tags) and preferences(user\_id, key, value)

**Schema (SQLite)**

```
CREATE TABLE preferences (user_id TEXT, key TEXT, value TEXT);
CREATE TABLE sessions (id TEXT PRIMARY KEY, user_id TEXT, query TEXT, summary TEXT, created_at DATETIME, metadata JSON);
```

Memory Agent exposes: `get_pref(user_id, key)`, `set_pref(user_id, key, value)`, `save_session(session_obj)`

## 8) Observability plan (logs/traces/metrics)

- **Logging:** Python `logging` module with structured JSON logs; log agent steps, tool calls, latencies.
- **Tracing:** Use OpenTelemetry to instrument key steps (research fetch, analysis summarization, critique evaluation).
- **Metrics:** Simple counters: `queries_processed`, `avg_research_time_ms`, `eval_score_mean`.

Example log line:

```
{"time": "2025-11-16T12:00:00+05:30", "agent": "research_agent",
"event": "search_called", "latency_ms": 123}
```

For Kaggle demo, logs can be printed to notebook and saved to `logs/` folder; for production, push to Cloud Logging.

## 9) Evaluation strategy (LLM-as-a-Judge + metrics)

**Metrics to report:** - **Coverage** (0-1): fraction of requested topics covered - **Citation quality** (0-1): are sources reliable? - **Clarity** (0-1): readability score (Flesch or LLM-based) - **Novelty/Redundancy** (0-1): duplication check against sources - **Overall score** = weighted sum

**LLM-as-a-Judge prompt template:**

```
You are an expert research reviewer. Evaluate the following student summary
on **clarity**, **accuracy**, **coverage**, **citation quality**, and
**originality**. Provide a numeric score (0-100) for each and a short
comment.
```

```
Summary:
"""
{summary}
"""
```

```
Sources:  
{sources}
```

Critique Agent returns JSON with scores; orchestrator maps to final grade.

---

## 10) Notebook & folder structure (Kaggle-ready)

```
IMRA-capstone/  
|   notebooks/  
|   |   └── imra_notebook.ipynb    # main notebook for demo  
|   src/  
|   |   agents/  
|   |   |   ├── orchestrator.py  
|   |   |   ├── research_agent.py  
|   |   |   ├── analysis_agent.py  
|   |   |   └── critique_agent.py  
|   |   tools/  
|   |   |   ├── web_search_tool.py  
|   |   |   ├── pdf_extractor.py  
|   |   |   └── chart_tool.py  
|   |   memory/  
|   |   |   └── memory_agent.py  
|   data/  
|   logs/  
|   requirements.txt  
└── slides/  
    └── presentation.pdf
```

**Kaggle notebook contains:** intro, quick config, run demo cells (3 example queries), visual outputs, evaluation table, downloadable artifacts.

---

## 11) Deployment plan (Vertex AI Agent Engine — optional but recommended)

**High-level steps** 1. Containerize agents as Python service (Flask/FastAPI). 2. Create Agent config for Vertex Agent Engine. 3. Setup authentication & IAM. 4. Deploy and test with sample inputs.

We'll include a `deploy/` folder with Dockerfile and a `vertex_deploy.md` with exact commands.

---

## 12) Demo script + testing plan

**Demo Input examples:** 1. "Summarize 3 papers on micro-Doppler signatures for human activity recognition" 2. "Compare Transformer vs CNN methods in 300 words"

**Expected outputs:** summary, 5 citations, critique with scores, two charts (keyword frequency, timeline of publication counts).

**Testing:** unit tests for tools (search returns list), pdf extractor accuracy, memory reads/writes, evaluation outputs JSON.

---

## 13) Slides / Presentation outline (10 slides)

1. Title + team + date
2. Motivation & problem statement
3. Key learnings leveraged (day-wise mapping)
4. High-level architecture (diagram)
5. Agent roles (one slide)
6. Demo screenshots + sample outputs
7. Observability & evaluation (metrics)
8. Deployment & scalability
9. Limitations & future work
10. Thank you + CTA (GitHub link)

Each slide will include succinct bullets and visuals.

---

## 14) LinkedIn capstone post (short + carousel caption)

**Main post:**

 I'm excited to share my capstone from the 5-Day AI Agents Intensive by Google & Kaggle: **IMRA — Intelligent Multi-Agent Research Assistant** (10–14 Nov 2025).

IMRA is a multi-agent system that finds, summarizes, critiques and visualizes research on demand using ADK, Gemini, A2A, MCP, sessions & memory, and observability. It returns a concise literature summary with citations, critique scores, and charts.

Demo & code: [GitHub link]

**Google #Kaggle #GeminiAI #ADK  
#VertexAI #AIAgents #MachineLearning  
#ResearchAssistant**

**Carousel captions:** 7 slides mapping to intro, architecture, demo, tech stack, observability, evaluation, thanks.

---

## 15) Scoring rubric alignment (how to maximize marks)

- Use >=6 elements (we use 8) — surpass requirement
  - Show live demo + screenshots (high impact)
  - Provide evaluation metrics & observability (technical depth)
  - Clear architecture diagrams + A2A flow (design quality)
  - Include optional deployment steps (production-readiness)
- 

## 16) Next steps & immediate deliverables (I will produce now if you want)

I can immediately deliver the following in this session:

1. Complete **Kaggle notebook** (`imra_notebook.ipynb`) with runnable demo cells (using mocked tools if external APIs blocked).
2. Full **src/** folder code (agents, tools, memory) as Python files.
3. **Slides** deck (10 slides) as exportable markdown or PDF-ready content.
4. **Presentation script** and **demo runbook**.
5. **GitHub-ready README** and `requirements.txt`.

Tell me if you want me to **generate the full Kaggle notebook and source code now**. If yes, I'll produce the complete code cells and files here in the notebook format ready to copy into Kaggle.

---

### Appendix: Helpful prompts

#### Orchestrator prompt (short):

You are Orchestrator. Given a user query, call Research Agent for sources, call Analysis Agent to summarize, call Critique Agent to evaluate, collect results and return final package with scores and citations.

#### Critique prompt (short):

Review the summary and score on Clarity, Accuracy, Coverage, and Citation Quality (0-100). Provide brief comments for improvement.

---

*End of document.*